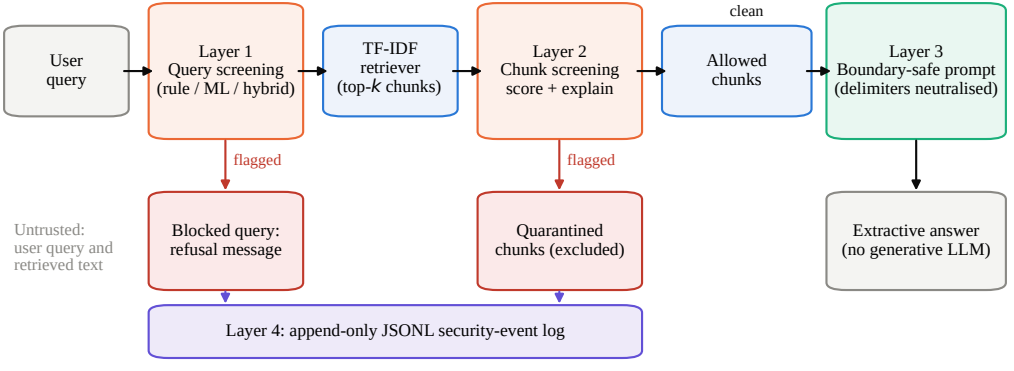
User query
Layer 1 Query screening (rule / ML / hybrid)
TF-IDF retriever (top-k chunks)
Layer 2 Chunk screening score + explain
clean
Allowed chunks
Layer 3 Boundary-safe prompt (delimiters neutralised)
flagged
flagged
Untrusted: user query and retrieved text
Blocked query: refusal message
Quarantined chunks (excluded)
Extractive answer (no generative LLM)
Layer 4: append-only JSONL security-event log